

DGL on Real World Applications

Quan Gan

AWS Shanghai AI Lab

April 19, 2020

Various GNN Business Applications

- Recommender Systems: find the most relevant items for a given user.
- Product Search: find the most relevant items given a set of keywords.
- Fraud Detection: detect anomaly of entities in a supervised/unsupervised manner.
- Community Detection: cluster entities in a supervised/unsupervised manner.

Recommender System: Problem Statement



Image source: Wikipedia

Collaborative Filtering



Collaborative Filtering




























Collaborative Filtering





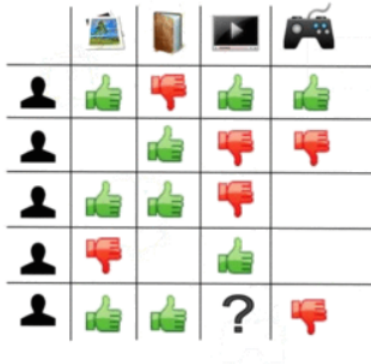
Image source: Wikipedia

Collaborative Filtering with Machine Learning

We were "representing" users and items with the items/users that had interactions with them.

Collaborative Filtering with Machine Learning



	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

=

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0?	3	0?	3	0?
User 2	4	0?	0?	2	0?
User 3	0?	0?	3	0?	0?
User 4	3	0?	4	0?	3
User 5	4	3	0?	4	0?

We were "representing" users and items with the items/users that had interactions with them.

Source: [Kat Bailey](#)

Can we represent users and items as a set of features?

Latent User/Item Representations

- An item can be described with a set of features (e.g. how sweet some food is).

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a set of features (e.g. how sweet some food is).
- A user can be described with preferences of the same set of features (e.g. how much a user likes sweet food).

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

=

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a set of features (e.g. how sweet some food is).
- A user can be described with preferences of the same set of features (e.g. how much a user likes sweet food).
- The interaction is defined by how well the item features match the user preferences.

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a **vector** v_j (sweet, organic, etc.).
- A user can be described with preferences of the same set of features (e.g. how much a user likes sweet food).
- The interaction is defined by how well the item features match the user preferences.

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with **another vector u_i** (likes sweet, likes organic, etc.)
- The interaction is defined by how well the item features match the user preferences.

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with another vector u_i (likes sweet, likes organic, etc.)
- The rating on item j by user i . is defined by $u_i^\top v_j$.

The diagram illustrates the process of matrix factorization. It shows a user-item rating matrix being decomposed into two lower-rank matrices, U and V , which are then multiplied back together to approximate the original matrix.

Original Matrix (User-Item Ratings):

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?
Feature 3	?	?	?	?	?
Feature 4	?	?	?	?	?
Feature 5	?	?	?	?	?

Decomposition:

The original matrix is decomposed into two matrices, U and V , which are then multiplied back together to approximate the original matrix.

Matrix U (User Latent Factors):

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Matrix V (Item Latent Factors):

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Approximation:

The original matrix is approximated by the product of U and V .

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with another vector u_i (likes sweet, likes organic, etc.)
- The rating on item j by user i . is defined by $u_i^\top v_j$.
- We minimize

$$\sum_{i,j} \left(r_{i,j} - \left(u_i^\top v_j \right) \right)^2$$

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: Kat Bailey

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with another vector u_i (likes sweet, likes organic, etc.)
- The rating on item j by user i . is defined by $u_i^\top v_j$.
- We minimize

$$\sum_{i,j} \left(r_{i,j} - \left(u_i^\top v_j + b_{u_i} + b_{v_j} \right) \right)^2$$

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with another vector u_i (likes sweet, likes organic, etc.)
- The rating on item j by user i . is defined by $u_i^\top v_j$.
- We minimize

$$\sum_{i,j} \left(r_{i,j} - \left(u_i^\top v_j + b_{u_i} + b_{v_j} \right) \right)^2$$

$+\alpha \mathcal{R}(U, V)$

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

=

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: Kat Bailey

Latent User/Item Representations

- An item can be described with a vector v_j (sweet, organic, etc.).
- A user can be described with another vector u_i (likes sweet, likes organic, etc.)
- The rating on item j by user i . is defined by $u_i^\top v_j$.
- We minimize

$$\sum_{(i,j) \in \mathcal{B}} \left(r_{i,j} - \left(u_i^\top v_j + b_{u_i} + b_{v_j} \right) \right)^2 + \alpha \mathcal{R}(U, V)$$

	Feature 1	Feature 2
User 1	?	?
User 2	?	?
User 3	?	?
User 4	?	?
User 5	?	?

X

	Item 1	Item 2	Item 3	Item 4	Item 5
Feature 1	?	?	?	?	?
Feature 2	?	?	?	?	?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	0.?	3	0.?	3	0.?
User 2	4	0.?	0.?	2	0.?
User 3	0.?	0.?	3	0.?	0.?
User 4	3	0.?	4	0.?	3
User 5	4	3	0.?	4	0.?

Source: [Kat Bailey](#)

What if we don't have ratings?



	Item 1	Item 2	Item 3	Item 4	Item 5
User 1		3	0		0
User 2	4			2	0
User 3	0			0	0
User 4		0	4		3
User 5	4	3	0	4	0

	i_1	i_2	i_3	i_4	
u_1	?	+	+	?	user
u_2	+	?	?	+	
u_3	+	+	?	?	
u_4	?	?	+	+	
u_5	?	?	+	?	
	item				

Source: BPR: Bayesian Personalized Ranking from

Implicit Feedback, Rendle et al. 2012

Implicit Feedback

- For a given user i , an item being interacted j should have a higher score than another item k which was never being interacted.

	i_1	i_2	i_3	i_4	
u_1	?	+	+	?	user ↑ ↓
u_2	+	?	?	+	
u_3	+	+	?	?	
u_4	?	?	+	+	
u_5	?	?	+	?	
	← item →				

Source: *BPR: Bayesian Personalized Ranking*
from *Implicit Feedback*, Rendle et al. 2012

Implicit Feedback

- For a given user i , an item being interacted j should have a higher score than another item k which was never being interacted.
- We maximize

$$\sum_{i,j,k \in I \setminus I_{u_i}} \log \text{sigmoid}((u_i^\top v_j - u_i^\top v_k))$$

	i_1	i_2	i_3	i_4	
u_1	?	+	+	?	user ↑ ↓
u_2	+	?	?	+	
u_3	+	+	?	?	
u_4	?	?	+	+	
u_5	?	?	+	?	
	← item →				

Source: *BPR: Bayesian Personalized Ranking*
from *Implicit Feedback*, Rendle et al. 2012

Implicit Feedback

- For a given user i , an item being interacted j should have a higher score than another item k which was never being interacted.
- We maximize


























$$\sum_{i,j,k \in I \setminus I_{u_i}} \log \text{sigmoid}((u_i^\top v_j - u_i^\top v_k))$$

- We usually *sample* one or multiple k when computing gradients (**negative sampling**).
 - Commonly uniformly, but adaptive sampling often helps.

	i_1	i_2	i_3	i_4	
u_1	?	+	+	?	user ↑ ↓
u_2	+	?	?	+	
u_3	+	+	?	?	
u_4	?	?	+	+	
u_5	?	?	+	?	
	← item →				


























Source: *BPR: Bayesian Personalized Ranking*
from *Implicit Feedback*, Rendle et al. 2012

Formulating Recommender Systems as Graph Learning


				
				
				
				
				
				

- Explicit Feedback
- Implicit Feedback

Formulating Recommender Systems as Graph Learning


				
				
				
				
				
				

- Explicit Feedback
- Implicit Feedback



- Edge Classification/Regression
- Link Prediction

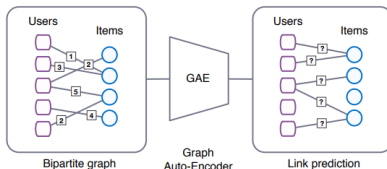
GCMC: Learning u_i and v_j from User-Item Graph



Source: *Graph Convolutional Matrix Completion*, van den Berg et al. 2017

1.
$$\mu_{v_j \rightarrow u_i, r} = \frac{1}{c_{u_i v_j}} W_r x_{v_j}$$


GCMC: Learning u_i and v_j from User-Item Graph



Source: *Graph Convolutional Matrix Completion*, van den Berg et al. 2017

1. $\mu_{v_j \rightarrow u_i, r} = \frac{1}{c_{u_i v_j}} W_r x_{v_j}$
2. $h_{u_i} = \sigma \left[\text{agg} \left(\sum_{v_j \in \mathcal{N}_{u_i, 1}} \mu_{v_j \rightarrow u_i, 1}, \dots, \sum_{v_j \in \mathcal{N}_{u_i, R}} \mu_{v_j \rightarrow u_i, R} \right) \right]$


GCMC: Learning u_i and v_j from User-Item Graph



Source: *Graph Convolutional Matrix Completion*, van den Berg et al. 2017

1. $\mu_{v_j \rightarrow u_i, r} = \frac{1}{c_{u_i v_j}} W_r x_{v_j}$
2. $h_{u_i} = \sigma \left[\text{agg} \left(\sum_{v_j \in \mathcal{N}_{u_i, 1}} \mu_{v_j \rightarrow u_i, 1}, \dots, \sum_{v_j \in \mathcal{N}_{u_i, R}} \mu_{v_j \rightarrow u_i, R} \right) \right]$
3. $u_i = \sigma(W_u h_{u_i})$ and similarly we compute v_j


GCMC: Learning u_i and v_j from User-Item Graph



Source: *Graph Convolutional Matrix Completion*, van den Berg et al. 2017

1. $\mu_{v_j \rightarrow u_i, r} = \frac{1}{c_{u_i v_j}} W_r x_{v_j}$
2. $h_{u_i} = \sigma \left[\text{agg} \left(\sum_{v_j \in \mathcal{N}_{u_i, 1}} \mu_{v_j \rightarrow u_i, 1}, \dots, \sum_{v_j \in \mathcal{N}_{u_i, R}} \mu_{v_j \rightarrow u_i, R} \right) \right]$
3. $u_i = \sigma(W_u h_{u_i})$ and similarly we compute v_j
4. $p(\hat{M}_{ij} = r) = \text{softmax}(u_i^\top Q_r v_j)$

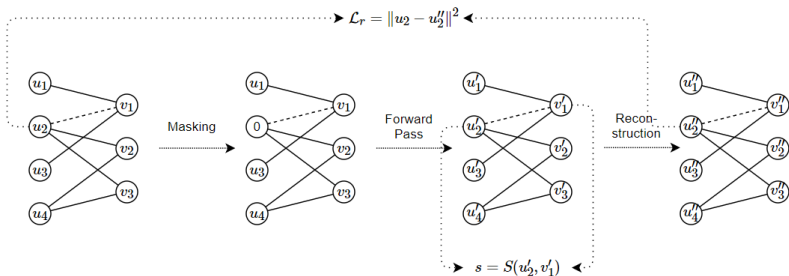
GCMC: Learning u_i and v_j from User-Item Graph



Source: *Graph Convolutional Matrix Completion*, van den Berg et al. 2017

1. $\mu_{v_j \rightarrow u_i, r} = \frac{1}{c_{u_i v_j}} W_r x_{v_j}$
2. $h_{u_i} = \sigma \left[\text{agg} \left(\sum_{v_j \in \mathcal{N}_{u_i, 1}} \mu_{v_j \rightarrow u_i, 1}, \dots, \sum_{v_j \in \mathcal{N}_{u_i, R}} \mu_{v_j \rightarrow u_i, R} \right) \right]$
3. $u_i = \sigma(W_u h_{u_i})$ and similarly we compute v_j
4. $p(\hat{M}_{ij} = r) = \text{softmax}(u_i^\top Q_r v_j)$
5. When new interactions are added, just re-run the forward pass on the new graph to get new u_i and v_j .


Learning u_i and v_j with Star-GCN



- Vanilla GCMC can't deal with new users/items without features (but with a few interactions).
- STAR-GCN
 - "Mask" the user/item embedding to 0 as if it is new.
 - Reconstruct the embedding after the forward pass and reconstruction pass.


Learning v_j and u_i from Item-Item and User-User Graph

- Decompose the user-item graph into user-user graph and item-item graph.




Learning v_j and u_i from Item-Item and User-User Graph

- Decompose the user-item graph into user-user graph and item-item graph.
- Get u_i with a Graph Convolutional Network on user-user graph and v_j on item-item graph.




Learning v_j and u_i from Item-Item and User-User Graph

- Decompose the user-item graph into user-user graph and item-item graph.
- Get u_i with a Graph Convolutional Network on user-user graph and v_j on item-item graph.
- Compute $r_{i,j} = u_i^\top v_j$




Learning v_j and u_i from Item-Item and User-User Graph

- Decompose the user-item graph into user-user graph and item-item graph.
- Get u_i with a Graph Convolutional Network on user-user graph and v_j on item-item graph.
- Compute $r_{i,j} = u_i^\top v_j$
- u_i and v_j can be learned with
 - Direct neighbor sampling (GraphSAGE)
 - Random-walk based neighbor sampling (PinSAGE)



Learning v_j and u_i from Item-Item and User-User Graph

- Decompose the user-item graph into user-user graph and item-item graph.
- Get u_i with a Graph Convolutional Network on user-user graph and v_j on item-item graph.
- Compute $r_{i,j} = u_i^\top v_j$
- u_i and v_j can be learned with
 - Direct neighbor sampling (GraphSAGE)
 - Random-walk based neighbor sampling (PinSAGE)
- When new interactions are added, just re-run the forward pass on the new graph to get new u_i and v_j .



Other Meaningful Aspects to Consider

- **Cold-start:** What if we have *new* users and items coming in, with few to no historical interactions?

Other Meaningful Aspects to Consider

- **Cold-start:** What if we have *new* users and items coming in, with few to no historical interactions?
- **Bias correction:** The training dataset usually comes from the result of a *previous recommender system*. How to mitigate the bias?

Other Meaningful Aspects to Consider

- **Cold-start:** What if we have *new* users and items coming in, with few to no historical interactions?
- **Bias correction:** The training dataset usually comes from the result of a *previous recommender system*. How to mitigate the bias?
- **Diversity:** Always recommending the same items (or even the same kind of item) to a user would make him/her feel *bored*.

Other Meaningful Aspects to Consider

- **Cold-start:** What if we have *new* users and items coming in, with few to no historical interactions?
- **Bias correction:** The training dataset usually comes from the result of a *previous recommender system*. How to mitigate the bias?
- **Diversity:** Always recommending the same items (or even the same kind of item) to a user would make him/her feel *bored*.
- **Fraud:** How to detect and deal with fabricated explicit feedbacks (e.g. fake ratings and reviews)?

Hands-on Session

Miniature GCMC on bipartite user-item graph.